

When Models Meet Data 2

Liang Zheng


Australian National University

liang.zheng@anu.edu.au

Overfitting

- The aim of a machine learning predictor is to perform well on **unseen data**.
- We simulate the unseen data by holding out a proportion of the whole dataset.
- This hold out set is called **test set**.
- In practice, we split data into a **training set** and a **test set**.
- **Training set**: fit the model
- **Test set**: not seen during training, used to evaluate generalization performance
- It is important for the user to not cycle back to a new round of training after having observed the test set.

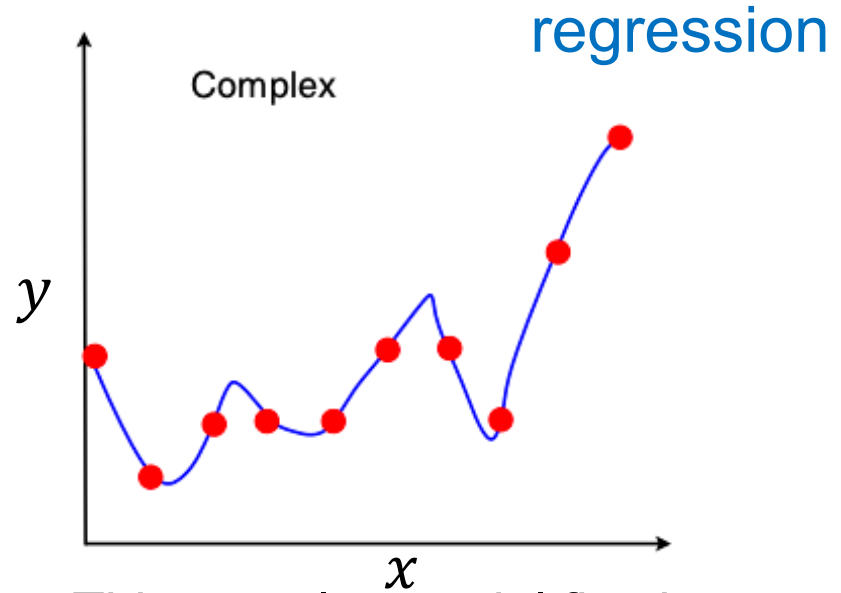
- Empirical risk minimization can lead to **overfitting**.
- the predictor fits too closely to the training data and does not generalize well to new data



This simple model fits the training data less well.

A larger empirical risk.

A good machine learning model.




This complex model fits the training data very well.

A very small empirical risk.

A poor machine learning model due to overfitting.


- Empirical risk minimization can lead to **overfitting**.
- the predictor fits too closely to the training data and does not generalize well to new data

A good model



● data, class 1
● data, class 2

A poor model **classification**




overfitted classification model

regularised classification model

8.2.3 Regularization to Reduce Overfitting

- When overfitting happens, we have
 - very **small** average loss on the training set but **large** average loss on the test set
- Given a predictor f , overfitting occurs when
 - the risk estimate from the training data $\mathbf{R}_{\text{emp}}(f, \mathbf{X}_{\text{train}}, \mathbf{y}_{\text{train}})$ underestimates the expected risk $\mathbf{R}_{\text{true}}(f)$. In other words,
 - $\mathbf{R}_{\text{emp}}(f, \mathbf{X}_{\text{train}}, \mathbf{y}_{\text{train}})$ is much smaller than $\mathbf{R}_{\text{true}}(f)$ which is estimated using $\mathbf{R}_{\text{emp}}(f, \mathbf{X}_{\text{test}}, \mathbf{y}_{\text{test}})$
- Overfitting occurs usually when
 - we have little data and a complex hypothesis class

- How to prevent overfitting?
- We can bias the search for the minimizer of empirical risk by introducing a penalty term
- The penalty term makes it harder for the optimizer to return an overly flexible predictor
- The penalty term is called **regularization**.
- Regularization is an approach that discourages complex or extreme solutions to an optimization problem.



- Example
- Least-squares problem

$$\min_{\boldsymbol{\theta}} \frac{1}{N} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|^2$$


- To regularize this formulation, we add a penalty term

$$\min_{\boldsymbol{\theta}} \frac{1}{N} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|^2 + \lambda \|\boldsymbol{\theta}\|^2$$

- The addition term $\|\boldsymbol{\theta}\|^2$ is called the **regularizer** or **penalty term**, and the parameter regularizer λ is the **regularization parameter**.
- λ enables a trade-off between **minimizing the loss on the training set** and the **amplitude of the parameters $\boldsymbol{\theta}$**
- It often happens that the **amplitude** of the parameters in $\boldsymbol{\theta}$ becomes relatively large if we run into overfitting
- λ is a hyperparameter


8.2.4 Cross-Validation to Assess the Generalization Performance

- We mentioned that we split a dataset into a **training set** and a **test set**
- we measure **generalization error** by applying the predictor on **test data**.
- This data is also sometimes referred to as the **validation set**.
- Validation set is from the entire data, and has no overlap with the training data.
- We want the training set to be **large**
- That leaves the validation set **small**
- A small validation set makes the **result less stable** (large variances)



- Basically, we want the training set to be large
- We want the validation to be large, too
- How to solve these contradictory objectives?
- **Cross-validation**: K -fold cross-validation


Example: $K = 5$



Cross-validation

- K -fold cross-validation partitions the data into K chunks
- $K - 1$ trunks form the training set \mathcal{R}
- The last trunk is the validation set \mathcal{V}
- This procedure is repeated for all K choices for the validation set, and the performance of the model from the K runs is averaged

Example: $K = 5$



Cross-validation

- Formally, we partition our training set into two sets $\mathcal{D} = \mathcal{R} \cup \mathcal{V}$, such that they do not overlap, i.e., $\mathcal{R} \cap \mathcal{V} = \phi$
- We train our model on \mathcal{R} (training set)
- We evaluate our model on \mathcal{V} (validation set)
- We have K partitions. In each partition k :
 - Training set $\mathcal{R}^{(k)}$ produces a predictor $f^{(k)}$
 - $f^{(k)}$ is applied to validation set $\mathcal{V}^{(k)}$ to compute the empirical risk $R(f^{(k)}, \mathcal{V}^{(k)})$
 - All the empirical risks are averaged to approximate the expected generalization error

$$\mathbb{E}_{\mathcal{V}}[R(f, \mathcal{V})] \approx \frac{1}{K} \sum_{k=1}^K R(f^{(k)}, \mathcal{V}^{(k)})$$

Cross-validation – some understandings

- The training set is limited -- not producing the best $f^{(k)}$
- The testing set is limited – producing an inaccurate estimation of $R(f^{(k)}, \mathcal{V}^{(k)})$
- After averaging, the results are stable and indicative
- An extreme: leave-one-out cross-validation, where the validation set only contains one example.
- A potential drawback – computation cost
 - The training can be time-consuming
 - If the model has several parameters to tune, it is hard to evaluate those hyperparameters.
- This problem can be solved by parallel computing, given enough computational resources

Check your understanding

- When your model works poorly on the training set, your model will also work poorly on the test set.
- When your model works poorly on the training set, your model may also have overfitting.
- Overfitting happens when your model is too complex given your training data.
- Regularization alleviates overfitting by improving the complexity of your training data.
- In K -fold cross-validation, we will get more stable test accuracy if K increases.
- In 2-fold cross-validation, you can obtain 2 results from the 2 test sets, and they may differ a lot with each other.