

Classification

Liang Zheng

Australian National University

liang.zheng@anu.edu.au

Reference: Bishop, Christopher M. *Pattern recognition and machine learning*. Springer, 2006.

Week 11 lecture arrangement

- Tuesday: 4pm-5:30pm, Copland lecture theatre & online (classification 1)
- Wednesday: 1pm-2:30pm, fully online (guest lecture)
- Thursday: 10:30am–12:00pm, fully online (classification 2)
- Friday: 3pm-4:30pm, Kambri cinema & online (review lecture)


For all these lectures, we use the same Zoom link for the online part:

<https://anu.zoom.us/j/85047332371?pwd=RDNOTnAvaG9VQnBDNzdGcnBINWpWZz09>

Meeting ID: 850 4733 2371

Password: 634639

Guest speakers




Liyue Shen, Postdoc@Harvard
University
Incoming Assistant Professor
@University of Michigan



José Lezama, Research Scientist
@Google Research

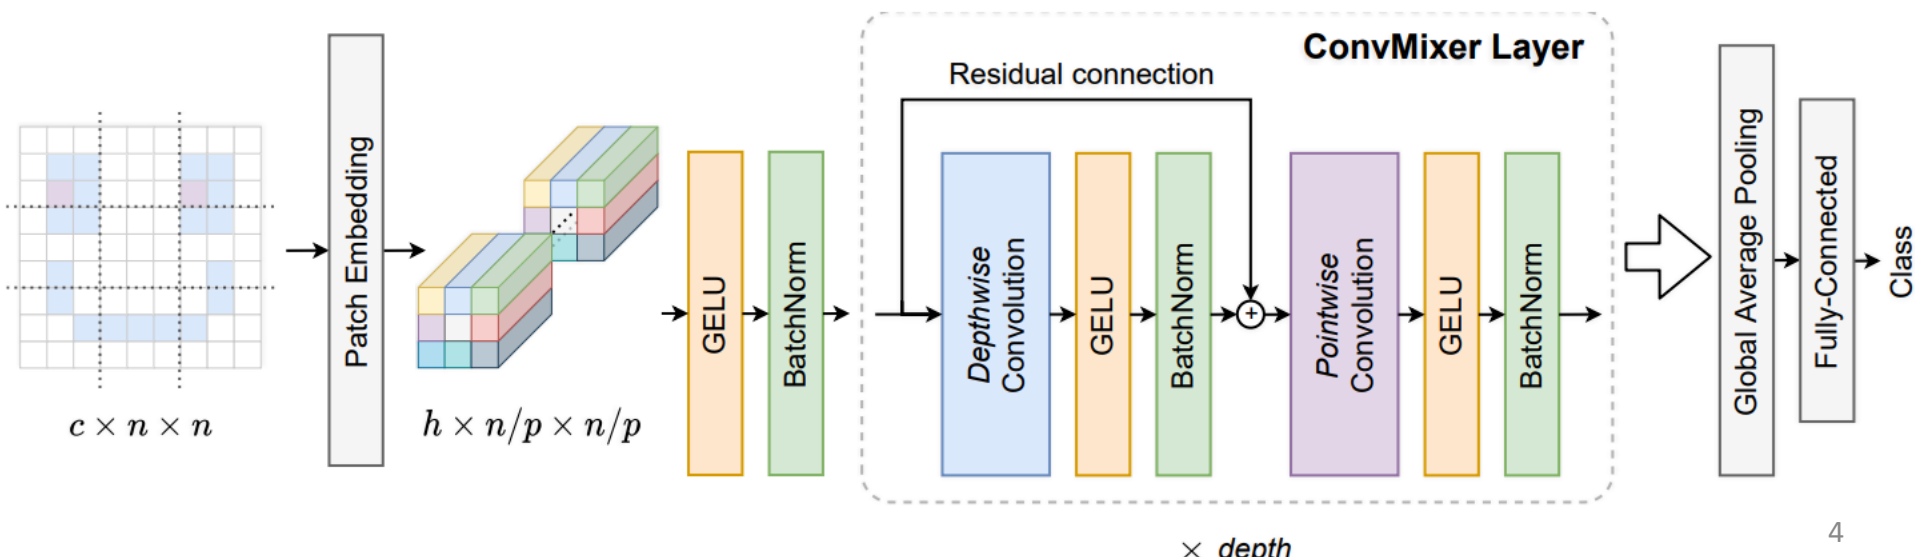
CONVOLUTIONS ATTENTION MLPs

PATCHES ARE ALL YOU NEED? 🙄




“A note on paper length. Expecting more text in this paper? Wondering if it’s a workshop paper we hastily submitted to ICLR? No. This paper presents a simple idea, one where we genuinely believe that a short paper presentation is more effective. Do we really need exactly 8 (now 9? 10?) pages to describe every machine learning architecture and algorithm in existence? We proposed an incredibly simple architecture and made a very simple point that we think is worth more discussion: patches work well in convolutional architectures. We think that four pages is more than enough space for this.”


★ ConvMixer ◆ ResMLP ● DeiT ■ ResNet



Classification problem



Goal: The goal in classification is to take an input vector \mathbf{x} and to assign it to one of K discrete classes $\mathcal{C}_k, k = 1, \dots, K$.



Goal: The goal in regression is to take an input vector \mathbf{x} and predict the value of a real number.

Regression

Training data

$$\mathcal{D} = \{ (\mathbf{x}_n, y_n) \mid n = 1, \dots, N \}$$

- Features/Inputs $\mathbf{x}_n \in \mathbb{R}^D$
- Response/Output $y_n \in \mathbb{R}$

Classification

Training data

$$\mathcal{D} = \{ (\mathbf{x}_n, y_n) \mid n = 1, \dots, N \}$$

- Features/Inputs $\mathbf{x}_n \in \mathbb{R}^D$
- Labels/Output $y_n \in \text{a set of discrete numbers}$

4.1.1 Two classes

- Linear classifiers $h: \mathbb{R}^D \rightarrow \{-1, +1\}$


$$h(\mathbf{x}; \boldsymbol{\theta}, \theta_0) = \text{sign}(y(\mathbf{x})) = \text{sign}(\boldsymbol{\theta}^T \mathbf{x} + \theta_0)$$

where $\boldsymbol{\theta}$ is called a **weight vector**, and θ_0 is a **bias** or **offset**. $y(\mathbf{x})$ is the **discriminant function**.

$$\text{sign}(y) = \begin{cases} +1 & \text{if } y \geq 0, \\ -1 & \text{if } y < 0. \end{cases}$$

An input vector \mathbf{x} is assigned to class $+1$ if $y \geq 0$ and to class -1 otherwise.

- The **decision boundary** is defined by the hyperplane $y(\mathbf{x}) = 0$, a $(D - 1)$ -dimensional hyperplane within the D -dimensional input space.




- Two points \mathbf{x}_A and \mathbf{x}_B both lie on the decision boundary $y(\mathbf{x}) = 0$. Then, we have $\boldsymbol{\theta}^T(\mathbf{x}_A - \mathbf{x}_B) = 0$ and hence the vector $\boldsymbol{\theta}$ is orthogonal to the decision boundary.
- Consider a point \mathbf{x} and let \mathbf{x}_\perp be its orthogonal projection onto the decision boundary. The perpendicular distance between \mathbf{x} and the decision boundary is given by,


$$r = \frac{|y(\mathbf{x})|}{\|\boldsymbol{\theta}\|}$$

- The distance from the origin to the decision surface is given by


$$\frac{|\boldsymbol{\theta}_0|}{\|\boldsymbol{\theta}\|}$$



Decision Regions



linear classifier




non-linear classifier

A classifier h partitions the space into **decision regions** that are separated by **decision boundaries**. In each region, all the points map to the same label. Many regions could have the same label.


4.1.1 Two classes

- Compact representation
- Similar to linear regression, it is convenient to use a more compact notation in which we introduce an additional dummy 'input' value $x_0 = 1$ and then define $\theta = [\theta_0, \theta_1, \theta_2, \dots, \theta_D]^T$ and $x = [x_0, x_1, x_2, \dots, x_D]^T$ so that
$$y(x) = \theta^T x$$
- In this case, the decision surfaces are D -dimensional hyperplanes passing through the origin of the $D + 1$ -dimensional expanded input space.
- The training data \mathcal{D} is **linearly separable** if there exist parameters $\theta = [\theta_0, \theta_1, \theta_2, \dots, \theta_D]^T$ such that for all $(x, y) \in \mathcal{D}$,
$$y(\theta^T x) > 0$$



4.1.2 Multiple classes

- Now we consider a linear classifier for $K > 2$ classes.



- One-versus-the-rest** classifier. Use $K - 1$ classifiers. Each classifier solves a two-class problem: separating points in a particular class C_k from points not in that class.
- One-versus-one** classifier. Use $K(K - 1)/2$ classifiers, one for every possible pair of classes. Each point is classified according to a majority vote amongst the classifiers.
- However, both methods suffer from the problem of ambiguous regions.

4.1.2 Multiple classes

- To avoid these difficulties, we consider a single K -class classifier comprising K linear functions of the form


$$y_k(\mathbf{x}) = \boldsymbol{\theta}_k^T \mathbf{x} + \theta_{k0}$$

- We then assign a point \mathbf{x} to class \mathcal{C}_k if $y_k(\mathbf{x}) > y_j(\mathbf{x})$ for all $j \neq k$.
- The decision boundary between class \mathcal{C}_k and class \mathcal{C}_j is therefore given by $y_k(\mathbf{x}) = y_j(\mathbf{x})$ and hence corresponds to a $(D - 1)$ -dimensional hyperplane defined by


$$(\boldsymbol{\theta}_k - \boldsymbol{\theta}_j)^T \mathbf{x} + (\theta_{k0} - \theta_{j0}) = 0$$

- This has the same form as the decision boundary for the **two-class** case discussed in Section 4.1.1, and so analogous geometrical properties apply.
- The decision regions of this classifier are singly connected and convex

“Singly connected and convex”: Two points \mathbf{x}_A and \mathbf{x}_B both lie inside decision region \mathcal{R}_k . Any point $\hat{\mathbf{x}}$ that lies on the line connecting \mathbf{x}_A and \mathbf{x}_B must also lie in \mathcal{R}_k




4.1.3 Least squares for classification



Our sample training dataset of 20 customers and the

You can use 0.5 as threshold to tell which class a data point belongs to



4.1.3 Least squares for classification

- We use the tools from least squares to solve classification.
- Problem setting
- Each class \mathcal{C}_k is described by its own compact linear model

$$y_k(\mathbf{x}) = \boldsymbol{\theta}_k^T \mathbf{x}$$

where $k = 1, \dots, K$. We group these together using vector notation

$$\mathbf{y}(\mathbf{x}) = \mathbf{\Theta}^T \mathbf{x}$$

- Where $\mathbf{\Theta}$ is a matrix whose k^{th} column comprises the $D + 1$ -dimensional vector $\boldsymbol{\theta}_k = (\theta_{k0}, \theta_{k1}, \dots, \theta_{kD})^T$ and $\mathbf{x} = (x_0, x_1, x_2, \dots, x_D)^T$ is the corresponding augmented input vector with a dummy input $x_0 = 1$.
- A new input \mathbf{x} is assigned to the class for which $y_k = \boldsymbol{\theta}_k^T \mathbf{x}$ is largest.

4.1.3 Least squares for classification

- We determine the parameter matrix Θ by minimizing the square error function
- Consider a training data set $\{\mathbf{x}_n, \mathbf{y}_n\}$ where $n = 1, \dots, N$, and define a matrix \mathbf{Y} whose n^{th} row is the vector \mathbf{y}_n^T , together with a matrix \mathbf{X} whose n^{th} row is \mathbf{x}_n^T . The square error function is written as

$$L(\Theta) = \frac{1}{2} \text{Tr}\{(\mathbf{X}\Theta - \mathbf{Y})^T(\mathbf{X}\Theta - \mathbf{Y})\}$$

- Setting the derivative with respect to Θ to zero, and rearranging, we obtain the closed-form solution for Θ :

$$\Theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y} = \mathbf{X}^\dagger \mathbf{Y}$$


- where \mathbf{X}^\dagger is the pseudo-inverse of \mathbf{X} . We obtain

$$\mathbf{y}(\mathbf{x}) = \Theta^T \mathbf{x} = \mathbf{Y}^T (\mathbf{X}^\dagger)^T \mathbf{x}$$

- $\mathbf{y}(\mathbf{x})$ is a $K \times 1$ -dim vector. The predicted class label corresponds to the largest value in $\mathbf{y}(\mathbf{x})$.

4.1.3 Least squares for classification


- Problems with Least squares



- If you still use 0.5 as the threshold, people between 20 and 22 will be misclassified.

4.1.3 Least squares for classification

- We can use Logistic regression to solve this problem
- Note: logistic regression is a classification method, not a linear regression method.



4.1.7 The Perceptron Algorithm

Rosenblatt (1962)

$\mathcal{L}_1(\boldsymbol{\theta}; \mathbf{x}, y)$ is the loss for one training sample (\mathbf{x}, y)

Classifier:

$$h(\mathbf{x}; \boldsymbol{\theta}) = \text{sign}(\boldsymbol{\theta}^T \mathbf{x})$$

Let $\mathcal{L}_1(\boldsymbol{\theta}; \mathbf{x}, y) = 1$ (0 otherwise) if

- $y \neq h(\mathbf{x}; \boldsymbol{\theta})$, or [misclassified]
- (\mathbf{x}, y) is on decision boundary [boundary]

Note that $y(\boldsymbol{\theta}^T \mathbf{x}) \leq 0$ if

- $\boldsymbol{\theta}^T \mathbf{x}$ and y differ in sign, or [misclassified]
- $\boldsymbol{\theta}^T \mathbf{x}$ is zero [boundary]

$$\mathcal{L}_1(\boldsymbol{\theta}; \mathbf{x}, y) = \mathbb{I}[y(\boldsymbol{\theta}^T \mathbf{x}) \leq 0] = \text{Loss}(y(\boldsymbol{\theta}^T \mathbf{x}))$$


where $\text{Loss}(z) = \mathbb{I}[z \leq 0]$ is the **zero-one loss**.

Training Loss

$$\text{Loss}(z) = \mathbb{I}[z \leq 0]$$

$$\mathcal{L}_1(\boldsymbol{\theta}; \mathbf{x}, y) = \text{Loss}\left(y(\boldsymbol{\theta}^T \mathbf{x})\right)$$

$$\mathcal{L}_n(\boldsymbol{\theta}; \mathcal{D}) = \frac{1}{n} \sum_{(\mathbf{x}, y) \in \mathcal{D}} \mathcal{L}_1(\boldsymbol{\theta}; \mathbf{x}, y)$$



Gradient is zero almost everywhere!

Gradient descent not possible.

Perceptron – a Mistake-Driven Algorithm

1. Initialize $\theta = 0$.
2. For each data $(x, y) \in \mathcal{D}$,
 - a. Check if $h(x; \theta) = y$.
 - b. If not, update θ to correct the mistake.
3. Repeat Step (2) until no mistakes are found.

Perceptron Algorithm

Weights may be initialized to $\mathbf{0}$ or to a small random value


1. Initialize $\theta = \mathbf{0}$.
2. For each data $(x, y) \in \mathcal{S}_n$,
 - a. If $y(\theta^\top x) \leq 0$,
 - i. $\theta \leftarrow \theta + yx$.
3. Repeat Step (2) until no mistakes are found.

Due to the constant feature trick $x_0 = 1$, update for θ_0 will be
 $\theta_0 \leftarrow \theta_0 + yx_0 = \theta_0 + y$.




Intuition behind $\theta \leftarrow \theta + yx$.

- We are making a mistake on a positive sample
- That is, $y = +1$, but $\theta^\top x \leq 0$
- According to this update, the new vector $\theta_{new} = \theta_{old} + yx = \theta_{old} + x$
- The new prediction will be
$$\theta_{new}^\top x = (\theta_{old} + x)^\top x = \theta_{old}^\top x + x^\top x > \theta_{old}^\top x$$
- For a positive example, the Perceptron update will increase the score assigned to the same input

Intuition behind $\theta \leftarrow \theta + yx$.



A mistake on a positive sample



Example

Training data

- $(\mathbf{x}_1, y_1) = ((2, 2)^\top, +1)$
- $(\mathbf{x}_2, y_2) = ((2, -1)^\top, -1)$



Apply the perceptron algorithm to the data to find a classifier $h(\mathbf{x}; \boldsymbol{\theta})$, $\boldsymbol{\theta} = (\theta_1, \theta_2)$, that separates the data.

Example

Training data

- $(\mathbf{x}_1, y_1) = ((2, 2)^\top, +1)$
- $(\mathbf{x}_2, y_2) = ((2, -1)^\top, -1)$

- Initialize $\boldsymbol{\theta} = (0, 0)^\top$.
- Since $y_1 \boldsymbol{\theta}^\top \mathbf{x}_1 = 0$,
set $\boldsymbol{\theta} = (0, 0)^\top + (2, 2)^\top = (2, 2)^\top$.
- Since $y_2 \boldsymbol{\theta}^\top \mathbf{x}_2 = -2$,
set $\boldsymbol{\theta} = (2, 2)^\top - (2, -1)^\top = (0, 3)^\top$.
- $y_1 \boldsymbol{\theta}^\top \mathbf{x}_1 = 6 > 0$.
- $y_2 \boldsymbol{\theta}^\top \mathbf{x}_2 = 3 > 0$.
- No more mistakes, so we are done.



Perceptron Algorithm

1. Training Set (Linearly Separable)

$$(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$$

2. Model (Set of Perceptrons)

$$h(\mathbf{x}; \boldsymbol{\theta}) = \text{sign}(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_D x_D)$$

3. Training Loss (Fraction of Misclassified/Boundary Points)

$$\mathcal{L}_n(\boldsymbol{\theta}) = \frac{1}{N} \sum_{(\mathbf{x}, y) \in \mathcal{D}} \mathbb{I}[y(\boldsymbol{\theta}^\top \mathbf{x}) \leq 0]$$


3. Algorithm (Mistake-Driven Algorithm)

Linearly Separable

The training data \mathcal{D} is
linearly separable
if there exists a
parameters θ and θ_0 such
that for all $(x, y) \in \mathcal{D}$,

$$y(\theta^\top x + \theta_0) > 0.$$





Perceptron algorithm can only be
applied to dataset that is linearly
separable



Check your understanding

- Classification deals with discrete label space.
- We can do multi-way classification using a single classifier.
- In multi-way classification $\mathbf{y}(\mathbf{x}) = \Theta^T \mathbf{x}$, the classifier weight θ_k can be shared by different classes.
- We can devise a classifier based on linear regression.
- In Perceptron algorithm, whenever a sample is misclassified, one step of θ update $\theta \leftarrow \theta + y\mathbf{x}$ will correct it.
- Trained on a linearly separable training set, a perceptron classifier will make no mistake on both training and test sets.

Data augmentation


	ResNet-50	Mixup [48]	Cutout [3]	CutMix
Image				
Label	Dog 1.0	Dog 0.5 Cat 0.5	Dog 1.0	Dog 0.6 Cat 0.4
ImageNet Cls (%)	76.3 (+0.0)	77.4 (+1.1)	77.1 (+0.8)	78.6 (+2.3)
ImageNet Loc (%)	46.3 (+0.0)	45.8 (-0.5)	46.7 (+0.4)	47.3 (+1.0)
Pascal VOC Det (mAP)	75.6 (+0.0)	73.9 (-1.7)	75.1 (-0.5)	76.7 (+1.1)

Zhang et al., mixup: Beyond Empirical Risk Minimization. ICLR 2018

Yun et al., CutMix: Regularization Strategy to Train Strong Classifiers with Localizable Features. ICCV 2019.

Zhong et al., Random Erasing Data Augmentation. AAAI 2020.


Hinge Loss



Perceptron algorithm does not converge for training sets that are not linearly separable.

Loss Functions

Training Loss



$$\mathcal{L}_n(\boldsymbol{\theta}) = \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \text{Loss}(y(\boldsymbol{\theta}^\top \mathbf{x})) \rightarrow z$$

Zero-One Loss

$$\text{Loss}_{01}(z) = \mathbb{I}[z \leq 0]$$

Hinge Loss

$$\text{Loss}_H(z) = \max\{1 - z, 0\}$$

CONVEX!


Penalize large mistakes more.

Penalize near-mistakes, i.e. $0 \leq z \leq 1$.

Hinge Loss

Find θ that minimizes

$$\begin{aligned}\mathcal{L}_n(\theta) &= \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \text{Loss}_H(z) \\ &= \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \max\{1 - z, 0\} \\ &= \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \max\{1 - y(\theta^\top x), 0\}\end{aligned}$$



Gradient

$$\nabla_z \text{Loss}_H(z) = \begin{cases} 0 & \text{if } z > 1, \\ -1 & \text{otherwise.} \end{cases}$$

$$\nabla_{\theta} \text{Loss}_H(y(\theta^\top x)) = \begin{cases} 0 & \text{if } y(\theta^\top x) > 1, \\ -yx & \text{otherwise.} \end{cases}$$

Stochastic Gradient Descent

1. Initialize $\boldsymbol{\theta} = \mathbf{0}$.
2. Select data $(\mathbf{x}, y) \in \mathcal{D}$ at random.
 - a. If $y(\boldsymbol{\theta}^\top \mathbf{x}) \leq 1$, then
 - i. $\boldsymbol{\theta} \longleftarrow \boldsymbol{\theta} + \eta_k y \mathbf{x}$.
3. Repeat Step (2) until convergence.
(e.g., when improvement in $\mathcal{L}(\boldsymbol{\theta})$ is small enough)

Differences from Perceptron Algorithm

- Check $z \leq 1$ rather than $z \leq 0$
- η_k rather than $\eta = 1$

Hinge Loss Algorithm

1. Training Set (Not Necessarily Linearly Separable)

$$(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$$

2. Model (Set of Perceptrons)

$$h(\mathbf{x}; \boldsymbol{\theta}) = \text{sign}(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_D x_D)$$

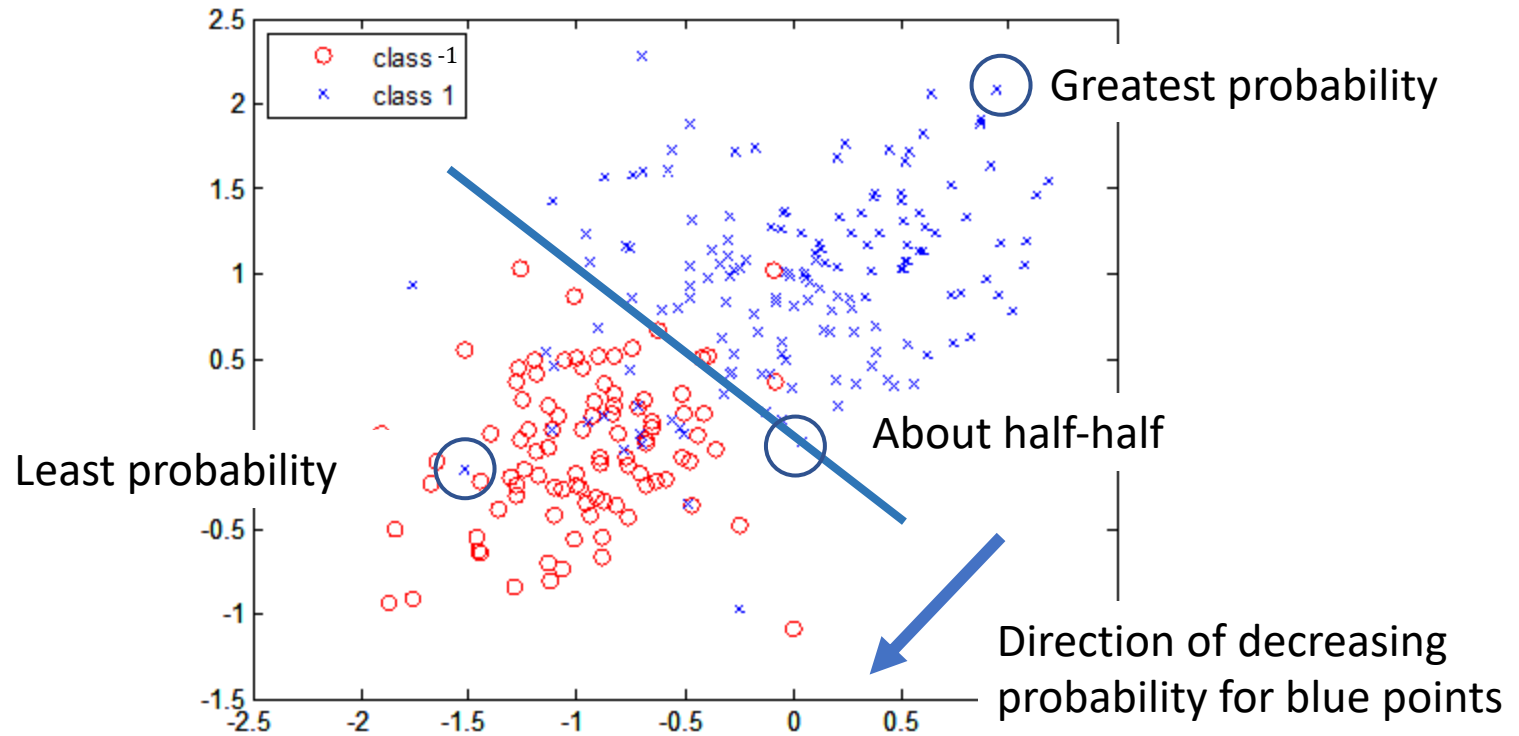
3. Training Loss (Hinge Loss)

$$\mathcal{L}_n(\boldsymbol{\theta}) = \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \max\{1 - y(\boldsymbol{\theta}^\top \mathbf{x}), 0\}$$

3. Algorithm (Gradient Descent)

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \frac{\eta_k}{N} \sum_{(x,y) \in \mathcal{D}} y \mathbf{x}$$

Logistic Regression




Probabilistic Model

Model the probability that the label y is $+1$ given the feature is \mathbf{x} .

$$h: \mathbb{R}^D \rightarrow [0, 1]$$

$$h(\mathbf{x}; \boldsymbol{\theta}) = \mathbb{P}(y = +1 | \mathbf{x}) = \text{sigmoid}(\boldsymbol{\theta}^\top \mathbf{x})$$



Sigmoid function


$$\text{sigmoid}(t) = \frac{1}{1 + e^{-t}}$$

$$\text{sigmoid}: \mathbb{R} \rightarrow [0, 1]$$

Sometimes also known as the **logistic** function.

Super useful formula

$$\text{sigmoid}(-t) = \frac{1}{1+e^t} = \frac{e^{-t}}{e^{-t}+1} = 1 - \frac{1}{e^{-t}+1} = 1 - \text{sigmoid}(t)$$




Label Probabilities

$$\mathbb{P}(y = +1 \mid \mathbf{x}) = \text{sigmoid}(\boldsymbol{\theta}^\top \mathbf{x}) = \text{sigmoid}(y(\boldsymbol{\theta}^\top \mathbf{x}))$$

To get the other label probability,

$$\begin{aligned}\mathbb{P}(y = -1 \mid \mathbf{x}) &= 1 - \mathbb{P}(y = +1 \mid \mathbf{x}) \\ &= 1 - \text{sigmoid}(\boldsymbol{\theta}^\top \mathbf{x}) \\ &= \text{sigmoid}(-\boldsymbol{\theta}^\top \mathbf{x}) = \text{sigmoid}(y(\boldsymbol{\theta}^\top \mathbf{x}))\end{aligned}$$

Thus, $\mathbb{P}(y \mid \mathbf{x}) = \text{sigmoid}(y(\boldsymbol{\theta}^\top \mathbf{x}))$ for both $y \in \{+1, -1\}$.




Label Predictions

Find out which label is more probable.

$$\mathbb{P}(y = +1 \mid \mathbf{x}) \geq \mathbb{P}(y = -1 \mid \mathbf{x}) \iff h(\mathbf{x}; \boldsymbol{\theta}) \geq \frac{1}{2}$$

If $h(\mathbf{x}; \boldsymbol{\theta}) \geq \frac{1}{2}$, then we predict the label of \mathbf{x} is $y = +1$.

If $h(\mathbf{x}; \boldsymbol{\theta}) < \frac{1}{2}$, then we predict the label of \mathbf{x} is $y = -1$.





Decision Boundary

$$h(\mathbf{x}; \boldsymbol{\theta}) \geq \frac{1}{2} \iff \text{sigmoid}(\boldsymbol{\theta}^\top \mathbf{x}) \geq \frac{1}{2} \iff \boldsymbol{\theta}^\top \mathbf{x} \geq 0$$

$$h(\mathbf{x}; \boldsymbol{\theta}) < \frac{1}{2} \iff \text{sigmoid}(\boldsymbol{\theta}^\top \mathbf{x}) < \frac{1}{2} \iff \boldsymbol{\theta}^\top \mathbf{x} < 0$$

The decision boundary
is described by $\boldsymbol{\theta}^\top \mathbf{x} = 0$.



Likelihood

Probability of label y given feature \mathbf{x}

$$\mathbb{P}(y|\mathbf{x}) = \text{sigmoid}(y(\boldsymbol{\theta}^\top \mathbf{x})).$$

$L(\boldsymbol{\theta}; \mathcal{D})$ is called the
likelihood of the dataset \mathcal{D} .

Probability of labels y_1, \dots, y_N given features $\mathbf{x}_1, \dots, \mathbf{x}_N$

$$\begin{aligned} L(\boldsymbol{\theta}; \mathcal{D}) &= \mathbb{P}(y_1, \dots, y_N | \mathbf{x}_1, \dots, \mathbf{x}_N) \\ &= \mathbb{P}(y_1 | \mathbf{x}_1) \times \dots \times \mathbb{P}(y_N | \mathbf{x}_N) \\ &= \prod_{(\mathbf{x}, y) \in \mathcal{D}} \mathbb{P}(y | \mathbf{x}) \end{aligned}$$

Maximizing $L(\boldsymbol{\theta}; \mathcal{D})$ is the same as **maximizing** $\log L(\boldsymbol{\theta}; \mathcal{D})$,
which is the same as **minimizing** $-\frac{1}{N} \log L(\boldsymbol{\theta}; \mathcal{D})$.

Logistic Loss

Minimize the training loss


$$\begin{aligned}\mathcal{L}(\boldsymbol{\theta}) &= -\frac{1}{N} \log L(\boldsymbol{\theta}; \mathcal{D}) \\ &= -\frac{1}{N} \log \prod_{(x,y) \in \mathcal{D}} \mathbb{P}(y|\mathbf{x}) \\ &= -\frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \log \mathbb{P}(y|\mathbf{x}) \\ &= -\frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \log \frac{1}{1+e^{-y(\boldsymbol{\theta}^\top \mathbf{x})}} \\ &= \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \log(1 + e^{-y(\boldsymbol{\theta}^\top \mathbf{x})}) \\ &= \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \text{Loss}(y(\boldsymbol{\theta}^\top \mathbf{x}))\end{aligned}$$

$\text{Loss}(z) = \log(1 + e^{-z})$
is the *logistic loss*.

Loss Functions

Training Loss

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \text{Loss}(y(\theta^\top x))$$



Zero-One Loss

$$\text{Loss}_{01}(z) = \mathbb{I}[z \leq 0]$$

$$z = y(\theta^\top x)$$

Perceptron Algorithm

Hinge Loss

$$\text{Loss}_H(z) = \max\{1 - z, 0\}$$

Hinge Loss Algorithm

Logistic Loss

$$\text{Loss}_L(z) = \log(1 + e^{-z})$$

Logistic Regression

Some folks use \log_2 instead of \log_e so that $\text{Loss}_L(0) = 1$.

Gradient

$$h(\mathbf{x}; \boldsymbol{\theta}) = \mathbb{P}(y = +1 | \mathbf{x}) = \text{sigmoid}(\boldsymbol{\theta}^\top \mathbf{x})$$

$$\text{sigmoid}(-t) = 1 - \text{sigmoid}(t)$$

$$\text{Loss}_L(z) = \log(1 + e^{-z})$$

$$\nabla_z \text{Loss}_L(z) = \frac{-e^{-z}}{1+e^{-z}} = \frac{-1}{e^z+1} = -\text{sigmoid}(-z) = \text{sigmoid}(z) - 1$$

By chain rule, the gradient for 1 training sample $\nabla_{\boldsymbol{\theta}} \mathcal{L}_1(\boldsymbol{\theta}; \mathbf{x}, y)$ is

$$\begin{aligned} \nabla_{\boldsymbol{\theta}} \text{Loss}_L(y(\boldsymbol{\theta}^\top \mathbf{x})) &= y\mathbf{x}(\text{sigmoid}(y\boldsymbol{\theta}^\top \mathbf{x}) - 1) \\ &= \begin{cases} \mathbf{x}(h(\mathbf{x}; \boldsymbol{\theta}) - 1) & \text{if } y = +1, \\ \mathbf{x}(h(\mathbf{x}; \boldsymbol{\theta}) - 0) & \text{if } y = -1. \end{cases} \\ &= \mathbf{x} (h(\mathbf{x}; \boldsymbol{\theta}) - \mathbb{I}[y = 1]) \end{aligned}$$

Training Gradient

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) &= \frac{1}{N} \sum_{(\mathbf{x}, y) \in \mathcal{D}} \nabla_{\boldsymbol{\theta}} \mathcal{L}_1(\boldsymbol{\theta}; \mathbf{x}, y) \\ &= \frac{1}{N} \sum_{(\mathbf{x}, y) \in \mathcal{D}} \mathbf{x} (h(\mathbf{x}; \boldsymbol{\theta}) - \mathbb{I}[y = 1])\end{aligned}$$

Compare this with the training gradient for least squares

$$\nabla_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) = \frac{1}{N} \sum_{(\mathbf{x}, y) \in \mathcal{D}} \mathbf{x} (f(\mathbf{x}; \boldsymbol{\theta}) - y)$$

In regression, we have:

$$\begin{aligned}L(\boldsymbol{\theta}, \mathbf{X}, \mathbf{y}) &= \frac{1}{N} \sum_{n=1}^N (y_n - f(\mathbf{x}_n, \boldsymbol{\theta}))^2 \\ f(\mathbf{x}_n, \boldsymbol{\theta}) &= \boldsymbol{\theta}^T \mathbf{x}_n\end{aligned}$$

Gradient Descent

1. Initialize θ randomly.
2. Update $\theta \leftarrow \theta - \frac{\eta_k}{N} \sum_{(x,y) \in \mathcal{D}} \mathbf{x}(h(\mathbf{x}; \theta) - \mathbb{I}[y = 1])$
3. Repeat (2) until convergence.
(e.g., when improvement in $\mathcal{L}(\theta; \mathcal{D})$ is small enough)

Logistic Regression

1. Training Set (Not Necessarily Linearly Separable)

$$(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$$

2. Model (Set of Sigmoid Neurons)

$$h(\mathbf{x}; \boldsymbol{\theta}) = \text{sigmoid}(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_D x_D)$$

3. Training Loss (Logistic Loss)

$$\mathcal{L}(\boldsymbol{\theta}) = \frac{1}{N} \sum_{(\mathbf{x}, y) \in \mathcal{D}} \log(1 + e^{-y(\boldsymbol{\theta}^\top \mathbf{x})})$$

3. Algorithm (Gradient Descent)

$$\boldsymbol{\theta} \longleftarrow \boldsymbol{\theta} - \frac{\eta_k}{N} \sum_{(\mathbf{x}, y) \in \mathcal{D}} \mathbf{x} (h(\mathbf{x}; \boldsymbol{\theta}) - \mathbb{I}[y = 1])$$

Regularized Logistic Regression


- When your data has a high-dimensional feature, or your training set is small, you might have the over-fitting problem.

$$\mathcal{L}(\boldsymbol{\theta}) = \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \log(1 + e^{-y(\boldsymbol{\theta}^\top x)}) + \frac{\lambda}{2N} \sum_{j=1}^D \theta_j^2$$


- We are doing regularization on $\theta_1, \theta_2, \dots, \theta_D$
- When using gradient descent, we have

$$\begin{aligned} \theta_0 &\leftarrow \theta_0 - \frac{\eta_k}{N} \sum_{(x,y) \in \mathcal{D}} x^{(0)} (h(\mathbf{x}; \boldsymbol{\theta}) - \mathbb{I}[y = 1]) \\ \theta_j &\leftarrow \theta_j - \frac{\eta_k}{N} \left[\sum_{(x,y) \in \mathcal{D}} x^{(j)} (h(\mathbf{x}; \boldsymbol{\theta}) - \mathbb{I}[y = 1]) + \lambda \theta_j \right] \\ &\hspace{15em} j = 1, 2, \dots, D \end{aligned}$$


Regularized Logistic Regression



overfitting



relatively good




underfitting

Lots of research in classification architecture / loss function design

- Softmax loss (softmax cross-entropy loss)
- **AM-Softmax** F. Wang, J. Cheng, W. Liu, and H. Liu. Additive margin softmax for face verification. IEEE Signal Processing Letters, 2018.
- **NormFace** F. Wang, X. Xiang, J. Cheng, and A. L. Yuille. Normface: L2 hypersphere embedding for face verification. ACM MM, 2017
- **Triplet loss with hard mining** F. Schroff, D. Kalenichenko, and J. Philbin. Facenet: A unified embedding for face recognition and clustering. CVPR 2015
- **Lifted structure loss** H. Oh Song, Y. Xiang, S. Jegelka, and S. Savarese. Deep metric learning via lifted structured feature embedding. CVPR 2016
- etc
- **Circle Loss** Sun, Yifan, et al. “Circle loss: A unified perspective of pair similarity optimization.” CVPR 2020.

Check your understanding



- Among the three loss functions, only the zero-one loss cannot be optimized by gradient descent.
- The logistic loss and hinge loss are generally superior to the zero-one loss.
- It is possible to sum the logistic loss and hinge loss as the final loss function, while keeping the system being able to be optimized by gradient descent.
- It is possible to sum the zero-one loss and hinge loss as the final loss function, while keeping the system optimized by gradient descent.
- “Easy” samples give very small loss values in these three loss functions.
- Comparing with the least square loss, the zero-one loss is more robust to “easy” samples.