

Review Lecture

Liang Zheng

Australian National University


liang.zheng@anu.edu.au

RepBin: Constraint-Based Graph Representation Learning for Metagenomic Binning



Hansheng Xue, Vijini Mallawaarachchi, Yujia Zhang, Vaibhav Rajan, Yu Lin. AAAI 2022

- An unsupervised approach to discovering clusters, of genomic subsequences, associated with the unknown constituent organisms


a) traditional contigs binning




b) RepBin




GroundTruth[n=10]



RepBin[n=10]



BusyBeeWeb[n=6]



CONCOCT[n=12]

Final exam

- **Time:** 9:30am – 13:00pm, Wednesday Nov 9th, 2021, AEDT.
- **Reading time:** 30 minutes.
- **Writing time:** 180 minutes
- **Zoom:** <https://anu.zoom.us/j/85047332371?pwd=RDNOTnAvaG9VQnBDNzdGcnBINWpWZz09>

Meeting ID: 850 4733 2371

Password: 634639

- **Total marks:** 100 (it weights 60% of your mark in this course)

Final exam

Some other details will be published on Piazza.

- You can bring any material (paper-based, electronic)
- You can use a physical calculator or your computer calculator. You can do any programming to verify your answer.
- You can do any internet search, any typing, but no communication devices / software
- You must record your screen throughout the exam
- Phones are not allowed during the exam. (You can use it to take photos of your answer after the exam ends).
- The exam paper will be published on Wattle at 9:30am.

Final exam

- **Exam scope:**
- All the lectures (Week 1 to 12).
- Math and machine learning will roughly weight 40 and 60, respectively.
- **Exam questions:**
- Proofs, calculations, derivations and short answers.

Final exam

- **Exam questions:**
- For short answer questions
 - If asked, you need to provide a short explanation to your answer (e.g., yes or no).
 - If your explanation is incorrect, you will get 0.
- Write down everything on paper.
- No programming, no pseudo codes.
- Exam questions from past years are on Wattle.
- The difficulty will be on a similar level with past exams (may be slightly more difficult or easier, subject to students' perception).
- Percentage of easy, medium and hard questions: 20:60:20 (roughly)
- **It is not a hurdle**

Final exam

- Keep your camera on all the time, so that I can see your face.
- Keep muted.
 - Type your question in the chat **privately to me**.
 - You must not type publicly.
 - The only ones who can type publicly is the Lecturer and tutors.
- You can freely touch your keyboard
 - typing a question privately.
 - Searching the internet
 - No communication software/device!
- Prepare plenty of papers to write your answers on
 - If you like, you can use iPad to write your answer, but you should make sure your iPad is screen-recorded

Final exam

- Write different sections on separate papers
- On your paper, please clearly denote the question numbers, e.g., 1(a), 2(b). You don't need to specify the section number, as your answers will be submitted to the corresponding section portal.
- Stop writing at 13:00pm. You will then have 20 minutes to upload your answers.
 - Take pictures of your answers.
 - On wattle, you will see the submission portal for different exam sections.
 - You can submit either 1 file or multiple files to each portal.
 - If you submit 1 file, please name it as 1.pdf, which is the compilation of all your answers.
 - If you submit multiple files (maximum 20), you should name them as 1.pdf, 2.pdf, 3.pdf,, according to the order of your answers.
 - You can also use other commonly seen files types, such as pdf, jpg, png.

Final exam

- Plagiarism in any form is absolutely forbidden.
- Probably fail the course if plagiarism is detected.
- There will be multiple versions of the exam questions
- According to our experience, because of the various question difficulties, some students couldn't finish all the exam questions. It is normal if you find yourself in a similar situation. We will make adjustment accordingly.
- Please report to me through email if you find anything inappropriate.
- Type in the chat privately if you want to go to the restroom

Linear algebra

- Matrix operations
- Systems of linear equations (e.g., Gaussian elimination)
- Vector subspaces
- Linear combination and linear independence
- Basis of a vector space, rank
- Linear mappings and matrix

Matrix decomposition

- Understand determinant and its properties
- Being able to calculate determinant
- Eigenvalues, eigenvectors, eigenspaces, eigendecomposition and SVD
- How eigendecomposition simplifies matrix operations?

Matrix decomposition

- Determinant (defined for square matrices)
- For 2×2 matrices, if $\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix}$, recall that the inverse of \mathbf{A} is

$$\mathbf{A}^{-1} = \frac{1}{a_{11}a_{22} - a_{12}a_{21}} \begin{bmatrix} a_{22} & -a_{12} \\ -a_{21} & a_{11} \end{bmatrix}$$

- For $n = 3$ (known as Sarrus' rule),

$$\begin{vmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{vmatrix} = a_{11}a_{22}a_{33} + a_{21}a_{32}a_{13} + a_{31}a_{12}a_{23} \\ - a_{31}a_{22}a_{13} - a_{11}a_{32}a_{23} - a_{21}a_{12}a_{33}$$

- For a square matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ it holds that \mathbf{A} is invertible if and only if $\det(\mathbf{A}) \neq 0$.
- Properties of determinant
- We can use Gaussian elimination to compute $\det(\mathbf{A})$ by bringing \mathbf{A} into row-echelon form. We can stop Gaussian elimination when we have \mathbf{A} in a triangular form where the elements below the diagonal are all 0. Recall: the determinant of a triangular matrix is the product of the diagonal elements.

Eigenvalues and Eigenvectors

- Let $A \in \mathbb{R}^{n \times n}$ be a square matrix. Then $\lambda \in \mathbb{R}$ is an **eigenvalue** of A and $x \in \mathbb{R}^n \setminus \{0\}$ is the corresponding **eigenvector** of A if

$$Ax = \lambda x$$

- We call this equation the **eigenvalue equation**.
- We calculate eigenvalues as the root of the **characteristic polynomial**

$$p_A(\lambda) := \det(A - \lambda I)$$

- Eigenspace
- The **eigenspace** of A with respect to λ and is denoted by E_λ
- Eigendecomposition
- A square matrix $A \in \mathbb{R}^{n \times n}$ can be factored into

$$A = PDP^{-1}$$

where $P \in \mathbb{R}^{n \times n}$ and D is a diagonal matrix whose diagonal entries are the eigenvalues of A , if and only if the eigenvectors of A form a basis of \mathbb{R}^n

Analytic geometry

- Understand norms, dot product and the more general inner product
- Lengths and distances are related to the inner product you choose
- Angles and orthogonality
- Linear mappings, eigenvalues, matrices, etc
- Orthogonal matrices $\mathbf{A} \mathbf{A}^T = \mathbf{I} = \mathbf{A}^T \mathbf{A}$ and rotations
- Orthonormal Basis
- Orthogonal projections and Gram-Schmidt Orthogonalization

Model meets data

- Understand model, predictor.
- Differentiate prediction, training and hyperparameter tuning.
- The role of cross validation
- Difference between parameters and hyperparameters
- Empirical risk minimization
- Evaluation metrics and loss functions
- The role of training data, testing data and validation data

Clustering

- Understanding EM algorithm for Clustering
- The difference between clustering and GMM
- Understanding the importance of key parameters in optimization.


Vector calculus

- Being able to calculate the gradient / partial derivatives of a function (with respect to vectors or matrices) and apply them to optimize machine learning problems.
- Being able to use identities in lectures without proofs.

Linear regression

- What is linear regression?
- The relationship between linear regression, PCA and classification
- Being capable of deriving the gradient and the closed-form solution of linear regression
- What is linear regression with features? How can the features help linear regression?

Linear regression



- Model:

$$y_n = \theta^T x_n$$

- Mean square error

$$\mathcal{L}(\theta) = \frac{1}{N} \|\mathbf{y} - \mathbf{X}\theta\|^2 = \frac{1}{N} (\mathbf{y} - \mathbf{X}\theta)^T (\mathbf{y} - \mathbf{X}\theta)$$

- \mathbf{X} is the data matrix; \mathbf{y} is the label vector; θ contains the parameters we want to optimize

Linear regression

- We calculate the gradient of \mathcal{L} with respect to the parameters $\boldsymbol{\theta}$ as


$$\begin{aligned}\frac{d\mathcal{L}}{d\boldsymbol{\theta}} &= \frac{d}{d\boldsymbol{\theta}} \left(\frac{1}{N} (\mathbf{y} - \mathbf{X}\boldsymbol{\theta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\theta}) \right) \\ &= \frac{1}{N} \frac{d}{d\boldsymbol{\theta}} (\mathbf{y}^T \mathbf{y} - \boldsymbol{\theta}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \boldsymbol{\theta} + \boldsymbol{\theta}^T \mathbf{X}^T \mathbf{X} \boldsymbol{\theta}) \\ &= \frac{1}{N} \frac{d}{d\boldsymbol{\theta}} (\mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X} \boldsymbol{\theta} + \boldsymbol{\theta}^T \mathbf{X}^T \mathbf{X} \boldsymbol{\theta}) \\ &= \frac{1}{N} (-2\mathbf{y}^T \mathbf{X} + 2\boldsymbol{\theta}^T \mathbf{X}^T \mathbf{X}) \in \mathbb{R}^{1 \times D}\end{aligned}$$

- The minimum is attained when the gradient is zero.

$$\begin{aligned}\frac{d\mathcal{L}}{d\boldsymbol{\theta}} &= \mathbf{0}^T \Leftrightarrow \boldsymbol{\theta}^T \mathbf{X}^T \mathbf{X} = \mathbf{y}^T \mathbf{X} \\ \Leftrightarrow \boldsymbol{\theta}^T &= \mathbf{y}^T \mathbf{X} (\mathbf{X}^T \mathbf{X})^{-1} \\ \Leftrightarrow \boldsymbol{\theta} &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}\end{aligned}$$

Linear regression

- Regularized linear regression
- Loss function



$$\mathcal{L}_\lambda(\boldsymbol{\theta}) = \frac{1}{N} \|\mathbf{y} - \Phi\boldsymbol{\theta}\|^2 + \lambda \|\boldsymbol{\theta}\|^2$$

- We calculate the gradient of \mathcal{L} with respect to the parameters $\boldsymbol{\theta}$ as

$$\begin{aligned} \frac{d\mathcal{L}}{d\boldsymbol{\theta}} &= \frac{d}{d\boldsymbol{\theta}} \left(\frac{1}{N} (\mathbf{y} - \mathbf{X}\boldsymbol{\theta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\theta}) + \lambda \|\boldsymbol{\theta}\|^2 \right) \\ &= \frac{1}{N} (-2\mathbf{y}^T \mathbf{X} + 2\boldsymbol{\theta}^T \mathbf{X}^T \mathbf{X}) + \frac{d}{d\boldsymbol{\theta}} (\lambda \|\boldsymbol{\theta}\|^2) \\ &= \frac{1}{N} (-2\mathbf{y}^T \mathbf{X} + 2\boldsymbol{\theta}^T \mathbf{X}^T \mathbf{X}) + 2\lambda \boldsymbol{\theta}^T \in \mathbb{R}^{1 \times D} \end{aligned}$$


- The minimum is attained when the gradient is zero.

$$\begin{aligned} \frac{d\mathcal{L}}{d\boldsymbol{\theta}} = \mathbf{0}^T &\Leftrightarrow 2\lambda \boldsymbol{\theta}^T + \frac{1}{N} 2\boldsymbol{\theta}^T \mathbf{X}^T \mathbf{X} = \frac{1}{N} 2\mathbf{y}^T \mathbf{X} \\ &\Leftrightarrow \boldsymbol{\theta}^T = \mathbf{y}^T \mathbf{X} (N\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X})^{-1} \\ &\Leftrightarrow \boldsymbol{\theta} = (N\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \end{aligned}$$


- Linear regression with features
- We are concerned with a regression problem $y = \phi^T(x)\theta + \epsilon$, where $x \in \mathbb{R}$ and $\theta \in \mathbb{R}^K$. An example transformation that is used in this context is

$$\phi(x) = \begin{bmatrix} \phi_0(x) \\ \phi_1(x) \\ \vdots \\ \phi_{K-1}(x) \end{bmatrix} = \begin{bmatrix} 1 \\ x \\ x^2 \\ x^3 \\ \vdots \\ x^{K-1} \end{bmatrix} \in \mathbb{R}^K$$


- We create a K -dimensional feature from a 1-dimensional input




(a) $M = 0$




(b) $M = 1$




(c) $M = 3$



(d) $M = 4$



(e) $M = 6$



(f) $M = 9$

Probability and distributions

- Being able to calculate discrete/continuous probability
- Being able to apply Bayes' Theorem
- Familiar with Gaussian distributions
- Understand mean, variance, expectation, covariance...

Probability and distributions

- Discrete probabilities
- Sample question:
- I have a bag, containing 5 balls of the same size. There are 3 white balls and 2 red balls. I randomly pick a ball from my bag. After each pick, the ball is not put back to the bag. If I make two picks.
- Q1: what is the probability that the two picked balls are both white?
- Q2: what is the probability that at least one of the two picked balls are white?

$$\text{Q1: } \frac{3}{5} \times \frac{2}{4} = \frac{3}{10}$$

$$\text{Q2: } 1 - \frac{2}{5} \times \frac{1}{4} = \frac{9}{10}$$

Sum Rule, Product Rule, and Bayes' Theorem

$$\underbrace{p(\mathbf{x} | \mathbf{y})}_{\text{posterior}} = \frac{\overbrace{p(\mathbf{y} | \mathbf{x})}^{\text{likelihood}} \overbrace{p(\mathbf{x})}^{\text{prior}}}{\underbrace{p(\mathbf{y})}_{\text{evidence}}}$$

Sample question

In the population, 5% of the males are color blind, while 0.25% of the females are color blind. Now we randomly select a person and find he/she is color blind.

Q: what is the probability that this person is male?

A: We use random variable M to represent male, F to represent female, C to represent color blindness.

We have

$$P(C|M) = 0.05, P(C|F) = 0.0025.$$

Using Bayes' theorem,


$$P(M|C) = \frac{P(M)P(C|M)}{P(M)P(C|M) + P(F)P(C|F)} = \frac{0.5 \times 0.05}{0.5 \times 0.05 + 0.5 \times 0.0025} = 95.2\%$$

Gaussian Mixture Models

- Understand GMMs, including
- Calculation with EM algorithm
- Understanding of the EM algorithm
- How to optimize the mean of the Gaussians

Gaussian Mixture Models

- A two-dimensional dataset



- A **Gaussian mixture model (GMM)** is a density model where we combine a finite number of K Gaussian distributions $\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$ so that

$$p(\mathbf{x}|\boldsymbol{\theta}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$
$$0 \leq \pi_k \leq 1, \sum_{k=1}^K \pi_k = 1$$

where we defined $\boldsymbol{\theta} := \{\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k, \pi_k: k = 1, \dots, K\}$ as the collection of all parameters of the GMM.

EM Algorithm

- Initialize μ_k, Σ_k, π_k . (below is an example)
 - $\pi_k = 1/K$ for all k
 - μ_k : centroids from k -means algorithm
 - $\Sigma_k = \Sigma$ the sample variance, for all k
- E-step: Evaluate responsibilities r_{nk} for every data point x_n using current parameters π_k, μ_k, Σ_k :

$$r_{nk} = \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}$$


- M-step: Re-estimate parameters π_k, μ_k, Σ_k using the current responsibilities r_{nk} (from E-step):

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N r_{nk} x_n \quad \leftarrow \text{You should be able to derive this}$$
$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N r_{nk} (x_n - \mu_k)(x_n - \mu_k)^T$$
$$\pi_k = \frac{N_k}{N}$$


Principal Component Analysis

- PCA's role in dimension reduction
- Understand the meaning of the intermediate data of PCA
- PCA calculation
- How PCA relates variance preservation with eigenvalues?
- PCA in high dimensions

Principal Component Analysis




(a) Dataset with x_1 and x_2 coordinates.




(b) Compressed dataset where only the x_1 coordinate is relevant.

PCA aims to find the direction where the variance is maximized.


Key Steps of PCA in Practice




(a) Original dataset.




(b) Step 1: Centering by subtracting the mean from each data point.




(c) Step 2: Dividing by the standard deviation to make the data unit free. Data has variance 1 along each axis.



(d) Step 3: Compute eigenvalues and eigenvectors (arrows) of the data covariance matrix (ellipse).



(e) Step 4: Project data onto the principal subspace.



(f) Undo the standardization and move projected data back into the original data space from (a).

eigenface

eigendecomposition

PCA in high dimensions

- The covariance matrix could be very large
- Instead calculating

$$\mathbf{S} = \frac{1}{N} \mathbf{X} \mathbf{X}^T \in \mathbb{R}^{D \times D}$$


- We calculate

$$\tilde{\mathbf{S}} = \frac{1}{N} \mathbf{X}^T \mathbf{X}$$

Classification

- Understand decision boundaries and decision regions
- Understand the differences between Perceptron, Hinge loss, and logistic loss
- Understand logistic regression
- Multi-way classification
- Differences and similarities between logistic regression and linear regression
- Logistic regression with regularization
- Being able to calculate gradient and apply gradient descent

Classification



linear classifier



non-linear classifier

A classifier h partitions the space into **decision regions** that are separated by **decision boundaries**. In each region, all the points map to the same label. Many regions could have the same label.

Classification

- The training data \mathcal{D} is **linearly separable** if there exist parameters $\theta = [\theta_0, \theta_1, \theta_2, \dots, \theta_D]^T$ such that for all $(x, y) \in \mathcal{D}$,
$$y(\theta^T x) > 0$$

Logistic regression vs linear regression



Perceptron

Rosenblatt (1962)

Classifier:

$$h(\mathbf{x}; \boldsymbol{\theta}) = \text{sign}(\boldsymbol{\theta}^T \mathbf{x})$$

Let $\mathcal{L}_1(\boldsymbol{\theta}; \mathbf{x}, y) = 1$ (0 otherwise) if

- $y \neq h(\mathbf{x}; \boldsymbol{\theta})$, or [misclassified]
- (\mathbf{x}, y) is on decision boundary [boundary]

Note that $y(\boldsymbol{\theta}^T \mathbf{x}) \leq 0$ if

- $\boldsymbol{\theta}^T \mathbf{x}$ and y differ in sign, or [misclassified]
- $\boldsymbol{\theta}^T \mathbf{x}$ is zero [boundary]

$$\mathcal{L}_1(\boldsymbol{\theta}; \mathbf{x}, y) = \mathbb{I}[y(\boldsymbol{\theta}^T \mathbf{x}) \leq 0] = \text{Loss}(y(\boldsymbol{\theta}^T \mathbf{x}))$$

where $\text{Loss}(z) = \mathbb{I}[z \leq 0]$ is the **zero-one loss**.


Logistic regression

- $\mathbb{P}(y = +1 \mid \mathbf{x}) = \text{sigmoid}(\boldsymbol{\theta}^\top \mathbf{x}) = \text{sigmoid}(y(\boldsymbol{\theta}^\top \mathbf{x}))$

- $\text{sigmoid}(t) = \frac{1}{1+e^{-t}}$

$$h(\mathbf{x}; \boldsymbol{\theta}) \geq \frac{1}{2} \iff \text{sigmoid}(\boldsymbol{\theta}^\top \mathbf{x}) \geq \frac{1}{2} \iff \boldsymbol{\theta}^\top \mathbf{x} \geq 0$$

$$h(\mathbf{x}; \boldsymbol{\theta}) < \frac{1}{2} \iff \text{sigmoid}(\boldsymbol{\theta}^\top \mathbf{x}) < \frac{1}{2} \iff \boldsymbol{\theta}^\top \mathbf{x} < 0$$



Logistic Loss

Minimize the training loss

$$\begin{aligned}\mathcal{L}(\boldsymbol{\theta}) &= -\frac{1}{N} \log L(\boldsymbol{\theta}; \mathcal{D}) \\&= -\frac{1}{N} \log \prod_{(x,y) \in \mathcal{D}} \mathbb{P}(y|\mathbf{x}) \\&= -\frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \log \mathbb{P}(y|\mathbf{x}) \\&= -\frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \log \frac{1}{1+e^{-y(\boldsymbol{\theta}^\top \mathbf{x})}} \\&= \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \log(1 + e^{-y(\boldsymbol{\theta}^\top \mathbf{x})}) \\&= \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \text{Loss}(y(\boldsymbol{\theta}^\top \mathbf{x}))\end{aligned}$$

$\text{Loss}(z) = \log(1 + e^{-z})$
is the *logistic loss*.

Regularized Logistic Regression

- When your data has a high-dimensional feature, or your training set is small, you might have the over-fitting problem.

$$\mathcal{L}(\boldsymbol{\theta}) = \frac{1}{N} \sum_{(x,y) \in \mathcal{D}} \log(1 + e^{-y(\boldsymbol{\theta}^\top x)}) + \frac{\lambda}{2N} \sum_{j=1}^D \theta_j^2$$

- We are doing regularization on $\theta_1, \theta_2, \dots, \theta_D$
- When using gradient descent, we have

$$\begin{aligned} \theta_0 &\leftarrow \theta_0 - \frac{\eta_k}{N} \sum_{(x,y) \in \mathcal{D}} x^{(0)} (h(\mathbf{x}; \boldsymbol{\theta}) - \mathbb{I}[y = 1]) \\ \theta_j &\leftarrow \theta_j - \frac{\eta_k}{N} \left[\sum_{(x,y) \in \mathcal{D}} x^{(j)} (h(\mathbf{x}; \boldsymbol{\theta}) - \mathbb{I}[y = 1]) + \lambda \theta_j \right] \\ &\hspace{15em} j = 1, 2, \dots, D \end{aligned}$$